

# **Innovating in a Software Graphics Pipeline**

**Paul Lalonde**

**Intel Corporation**

# User innovation matters

## Traditional API

- Mostly fixed function
- Some programmable points

## Extended Pipeline

- More fixed function
- More programmable parts
- But still under vendor control

## Extensible Pipeline

- Usable as-is
- Data structure & process conventions
- User can open it up and extend it

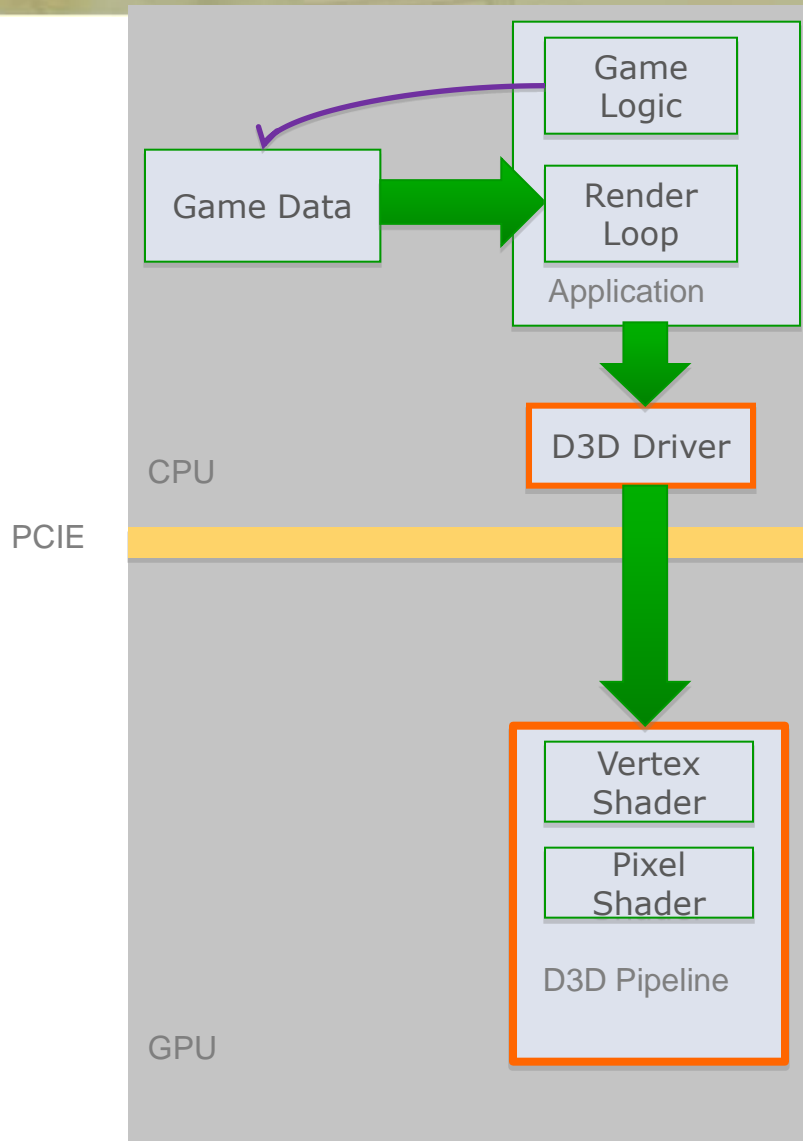
## Bare metal

- Totally flexible
- Only feasible to a few elite devs

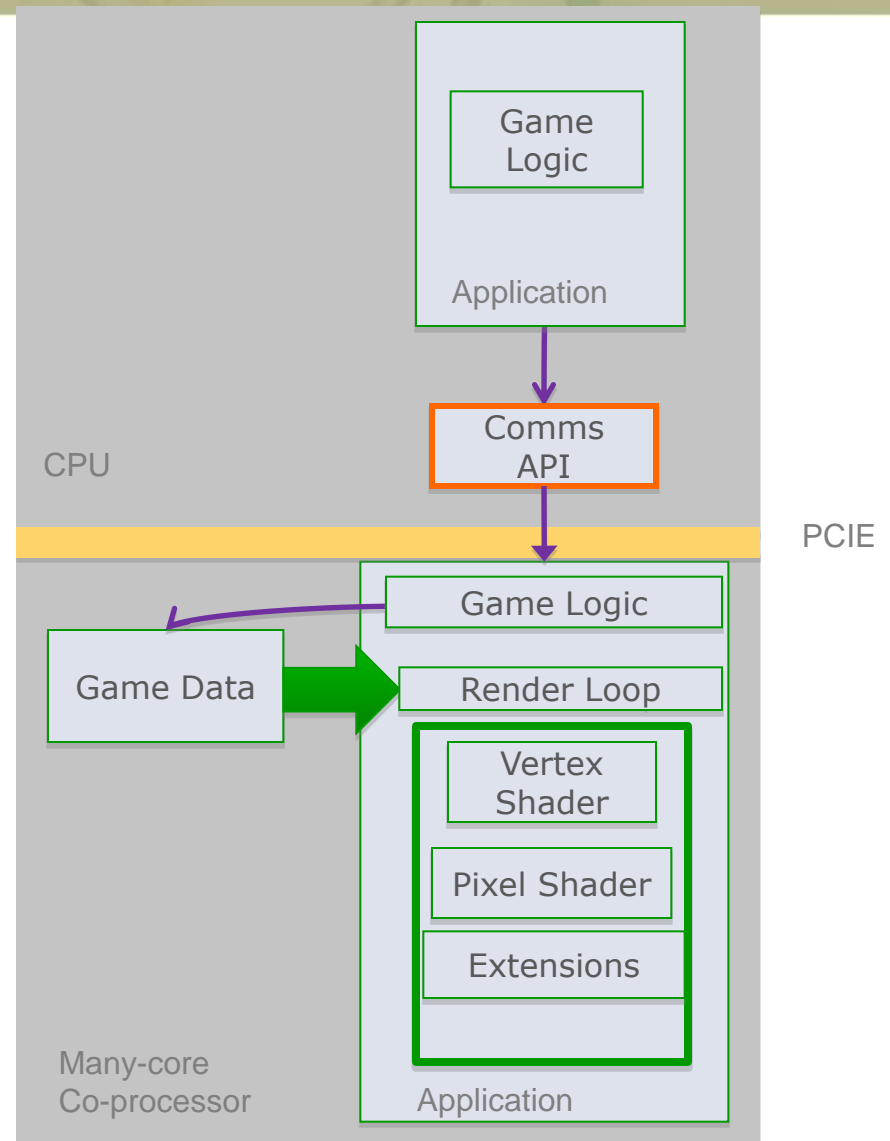
# Graphics APIs

- State-based, programmable at particular stages
- Restricted data model
- Innovation bottlenecks
  - Years of delay from proposal of feature to acceptance in the standard
  - Extensions either not supported (DX), or supported in ad hoc, vendor-incompatible ways (OpenGL)
- To be fair, also the best existing parallel programming model because of many of these restrictions

## Traditional Rendering Model



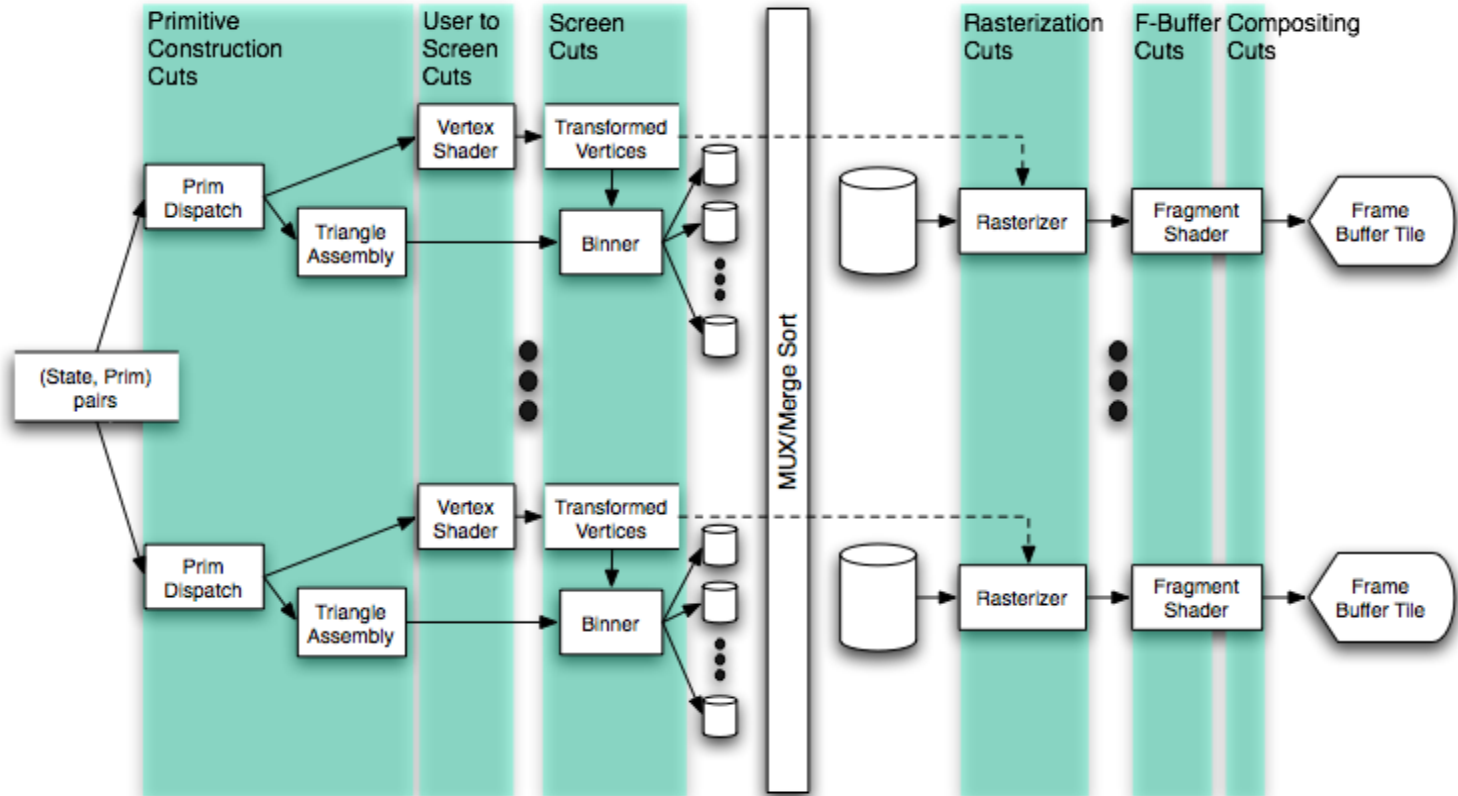
## Software Rendering Model



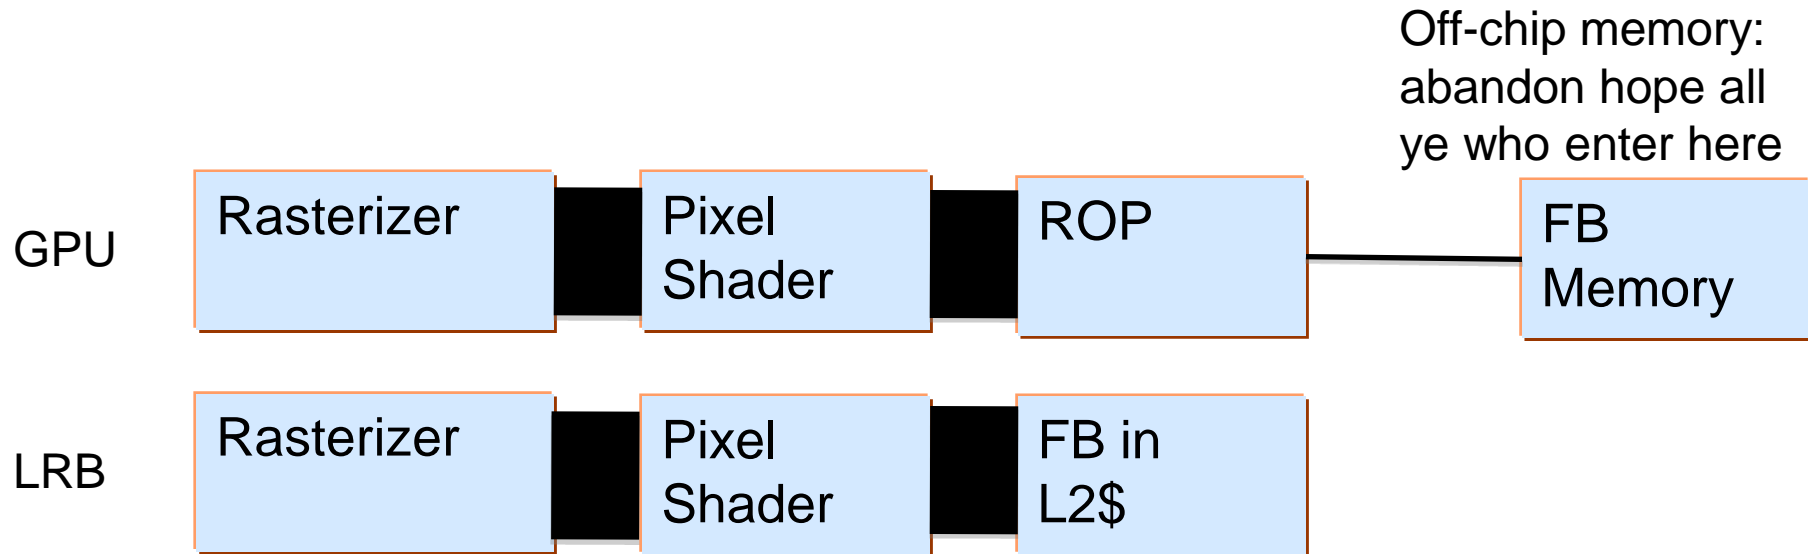
# Many-Core Graphics

- C/C++/Parallel Languages and Libraries
  - Throughput constraints
    - VPU utilization is important, “shader compilers” are important, even in user space (See Tim Foley’s talk)
- WDDM complements our memory model
  - Heap, resources can be declared on host
  - WDDM ensures declared resources are resident when LRB Native code is called
- Pushbuffer model hides PCIe latency
- Frame buffer presented without tunneling back to host

# A Rendering Organization



# Framebuffer RMW



- Pixel shader can read from current render target
- Compute new value based on existing one
- Blending can all happen with simple math in pixel shader
- More generally, pixel shader can maintain and update a small data structure at each pixel

# Read-modify-write


```
void myPixelShader(  
    ...  
    in out float fbDepth : SV_Depth,  
    in out float4 fbValue : SV_Target)  
{  
    // depth test in user code  
    if (fragPos.z > fbDepth)  
        return;  
    fbDepth = fragPos.z;  
  
    // “blend” using matrix multiply  
    fbValue = mul(m, fbValue);  
}
```

In out parameter  
bound to framebuffer  
semantic

Shader can read  
previous values from  
color/depth buffers at  
current pixel




# Translucency



You must sort by depth




So you have to sort per fragment



# K-Buffers (Bavoil et al. 2007)

- Store  $k$  frame buffers containing layer data
- Layered depth images, layered shadows
- (Limited) order independent transparency
  - Sort  $k$  layers at once, but can't sort any more than that
- Fixed allocation— $k$  layers
- Performance vs. flexibility trade-off
- Falls out easily from frame buffer RMW
  - Allocate  $n$  render targets
  - Do sorting/ordering in shader

# K-Buffer Transparency



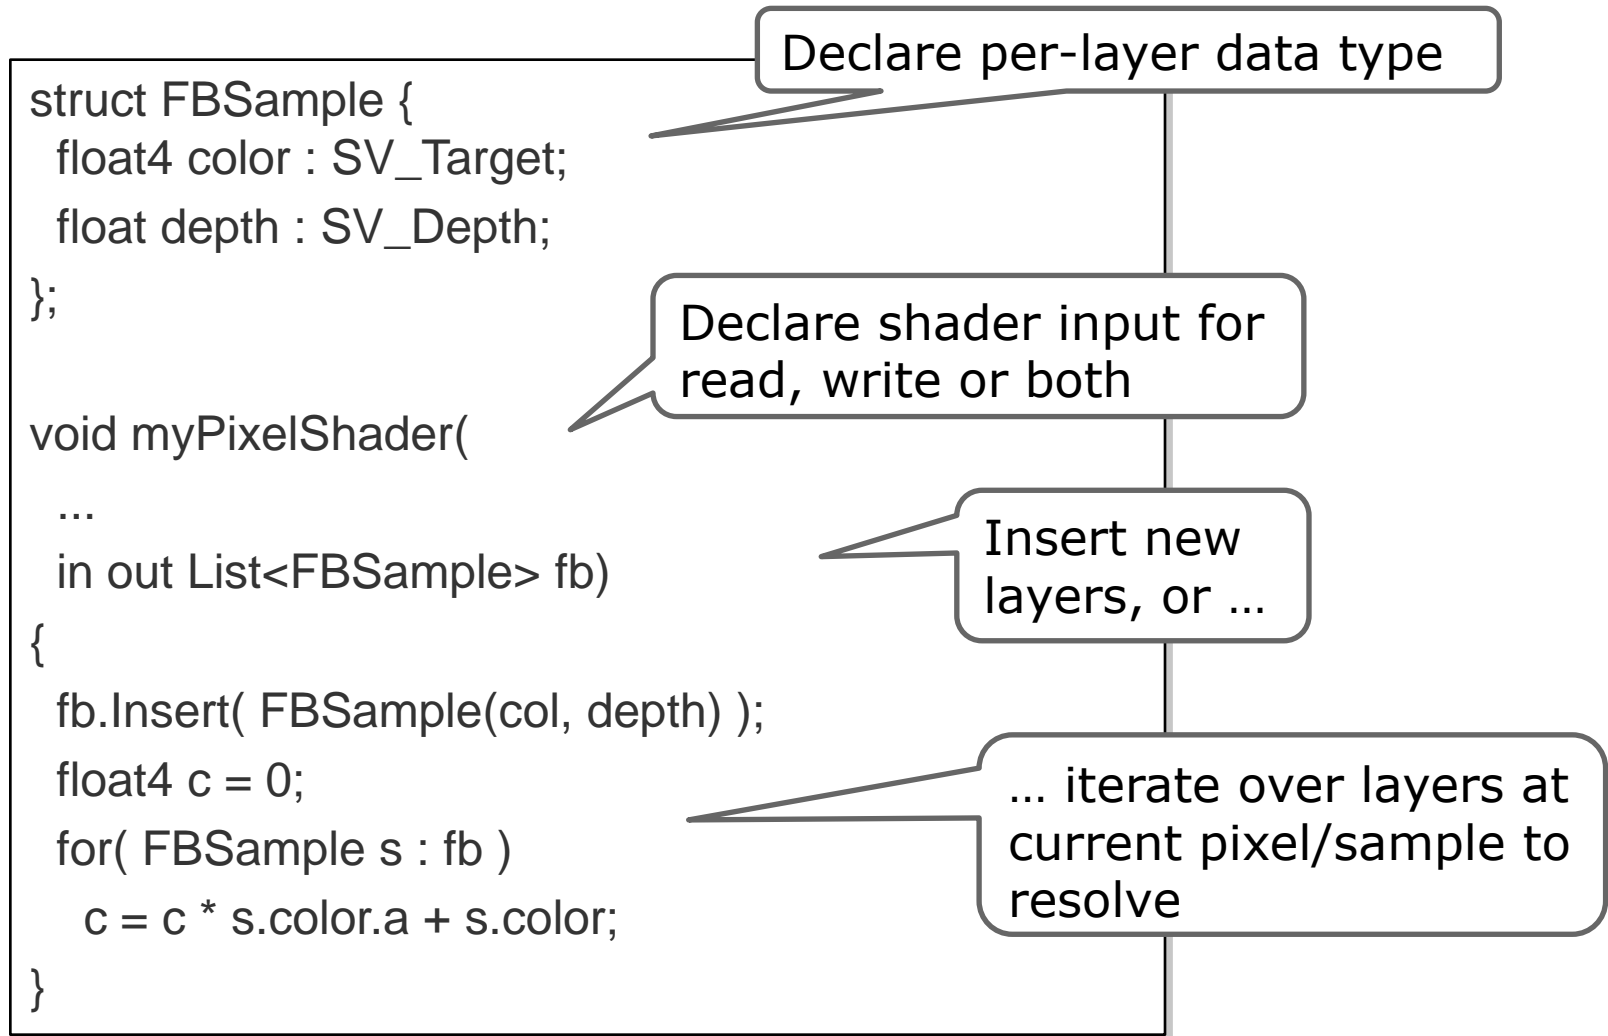
# List Render Targets (LRT)

- Allow arbitrary number of layers to be stored per pixel
  - Can store *all* fragments in a pixel, not just closest
  - Do in “user space”: let shader writer decide which ones to store, when to remove old ones, etc
- Render objects as usual
  - List of fragments is accumulated at each pixel
- Render quad running resolve shader
  - Sort if needed, composite layers, compute output values
  - Run out of L2\$, avoid main memory bandwidth


# List Textures

- Naïve implementation is linked list per pixel
- Smarter is linked list per SIMD width cluster of fragments, with masks
- Lists can be unsorted, sorted on insertion, post-sorted
- Principal optimization need is cache residence of working set of fragments & memory management within the List Texture
  - Thread safety of insertions & sorts
  - Smart cache management within a tile's working set
- Exposed as a List<> templated type in the shader


# List Render Targets



# LRT Examples








# Programmable Resolve

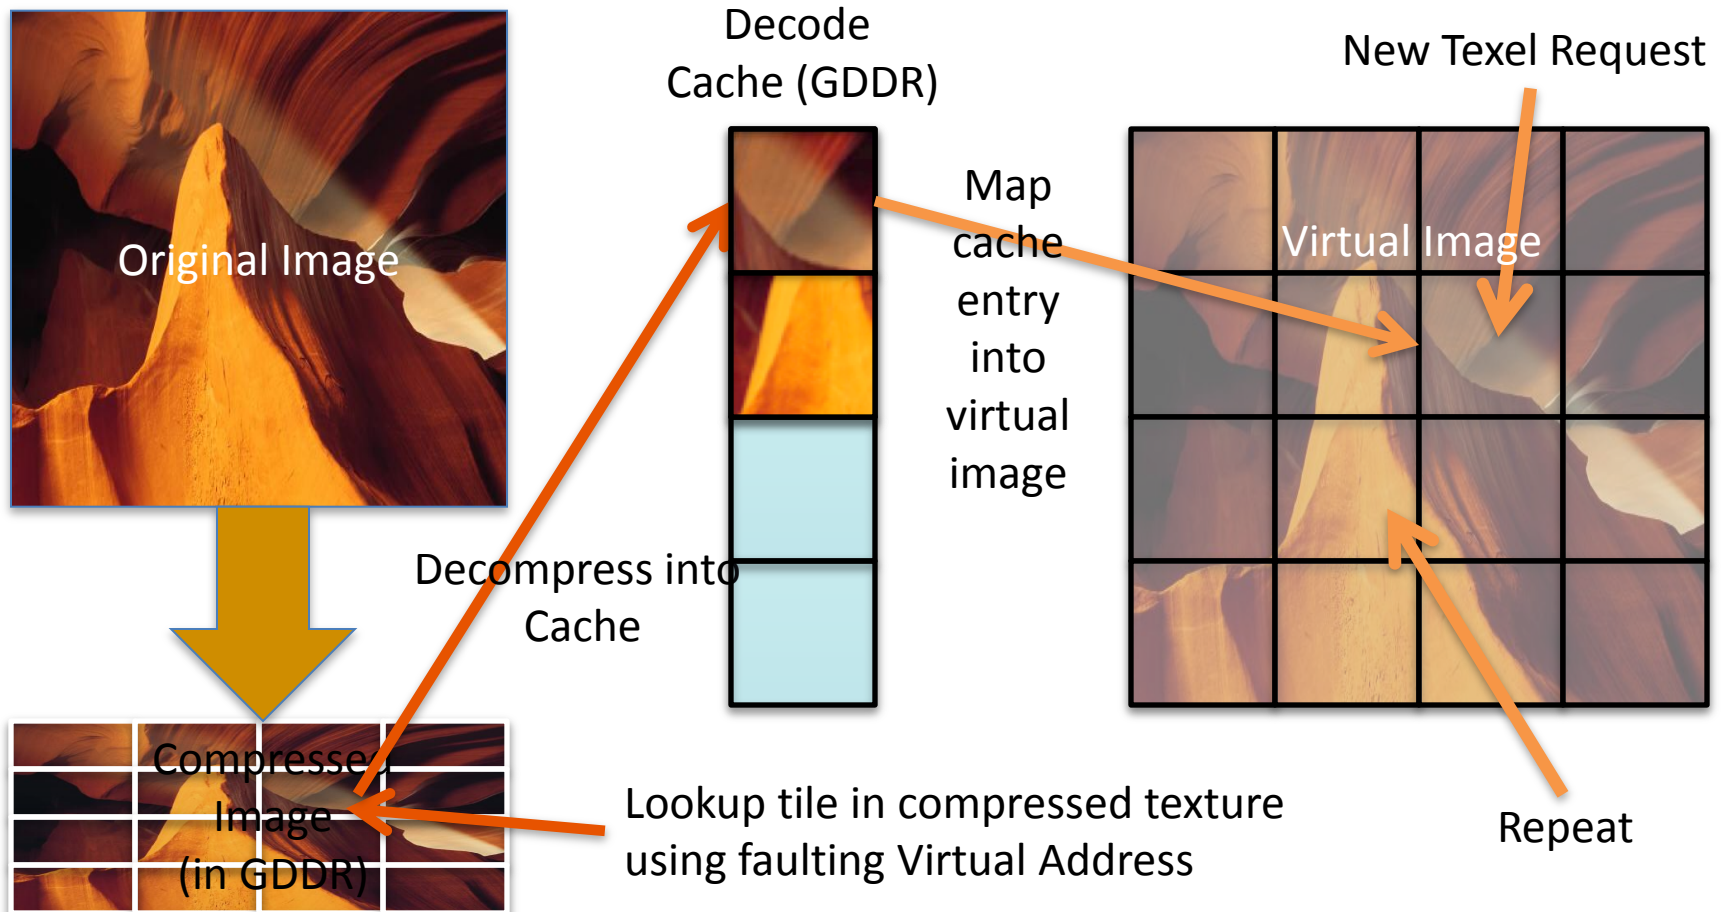
- After rasterizer backend finishes a tile, framebuffer is still in L2\$
- Useful opportunity to do image processing before write to memory
- Typical GPU approach:
  - Render to texture (commit to main memory)
  - Draw full-screen quad, read from texture (read back from memory)
  - Bandwidth intensive
  - Limited ability to leverage caches for reuse

# Programmable Resolve

- Call out to user-space code with pointer to framebuffer tile
  - Do all processing out of L2\$
  - No off-chip bandwidth cost
  - Uses: image processing, building data structures, reductions, ...
- Can do processing either with data-parallel code or with just a few threads
  - High performance available from either
  - Choose one where computation is more easily expressed



# Demand-Paged Textures



# Conclusion

- Each of these techniques is more expensive than straight rasterize/shade
- But each can save multiple geometry submission passes and GPU->CPU->GPU traffic, leading to higher system efficiency
- None of these extensions requires particular GPU hardware features from the vendor.
  - All can be done in user space
  - The value of the software pipeline is that it is extensible, not just extended.

Larrabee allows the ultimate  
in extensible pipelines

# Thank you

# Backup



# Shader Support for DPT

- Generated code looks like:

```
while(!done) {  
    err = sample(texture, u, v)  
    if(err == PNP) {  
        UserCB(cpuPage, txsPage, tileNo, userContext)  
    } else {  
        done = true  
    }  
}
```

# Managing the DPT Cache

- Many update strategies are possible
  - Just in time, Mip-Level surrogates, just-too-late
- Use any available physical pages to hold computed/decoded textures
- Allow the paging daemon to discard these instead of swapping them.
- The OS knows pretty well what's in use and what's stale, so why duplicate that in the DPT implementation?
- Multiple caches can be managed independently for different uses: worlds vs characters, for example

# Execution Model

- The renderer buffers up an entire frame's worth of draw commands
- This makes it possible for it to do **all** front-end graphics processing before starting back-end
- *Thus, changing state and/or switching render targets is cheap*
- *Mostly pointer updates in the same memory space*
- Also allows a nice alternative for submission
- Typical GPU:
  - For each render target, draw each visible object
- Software Rendering:
  - For each object, draw to all relevant render targets